TRmorph: A morphological analyzer for Turkish

Çağrı Çöltekin

June 19, 2015

1 Introduction

This document describes TRmorph version 0.9.

TRmorph is an open-source[1] finite-state morphological analyzer for Turkish. This document describes how to use the tools that comes with this package, as well as some implementation details that may be helpful for people who want to customize this open-source tool for their own needs. The source code is available at `https://github.com/coltekin/TRmorph`, and of the analyzer and a web-based demo can be accessed at `http://coltekin.net/cagri/trmorph`.

This document describes the current version of the analyzer and additional tools described in Çöltekin (2014). This version is a complete rewrite of the earlier version report in Çöltekin (2010). The earlier version of TRmorph was implemented using SFST (Schmid 2005), the current version is implemented with more popular finite state description languages *lexc* and *xfst* from Xerox (Beesley and Karttunen 2003), using Foma (Hulden 2009) as the main development tool.

The lexc/xfst implementation of TRmorph should compile with any lexc/xfst compiler without much additional effort. The only foma-specific notation used in the morphology description is about handling simple reduplication, which can also be handled with twolc rules, or compile-replace (Beesley and Karttunen 2003).[2]

2 How to use it

2.1 Compilation from the source

To compile TRmorph from the source, you need a lexc/xfst compiler such as foma, a C preprocessor, GNU make and some standard UNIX utilities.

If all requirements are in place, to build analyzer/generator FST, you should type `make` in the main TRmorph distribution directory. The resulting binary automaton file will be `trmorph.fst`.

Trmorph comes with a set of other finite-state tools that are useful in various NLP tasks. Currently, tolls for the following tasks are distributed together with TRmorph.

- stemming/lemmatization
- morphological segmentation
- hyphenation
- guessing unknown words

To compile these tools, you should specify the FST you want to build as an argument to `make`, e.g., `make stemmer` will build an binary automaton called `stem.fst`. These additional tools are described in Section 6.

2.2 Customizing TRmorph

TRmorph is an open source utility. As a result, you are free to modify the source according to your needs. Source code includes some useful comments on what/how/where things are done. Furthermore, TRmorph can be customized for some common choices during the compilation. These options are typically related to more relaxed analysis. For example whether to allow non-capitalized proper names, or analyze (and generate) text written in all

---

[1] Current version of TRmorph is licensed under GNU Lesser General Public License. See README file in the TRmorph distribution for more information.

[2] TRmorph can be compiled with HFST (Lindén et al. n.d.) without modification since HFST uses foma as the back end for parsing xfst files.

capitals, or set the decimal and thousand separator in numbers. These options are set in the file `options.h`. The file contains documentation along with the existing options. This feature is under development (as of July 2013), new options are currently being added, and existing options may not fully work as intended yet.

Another common need for customizing a morphological analyzer is to add or modify the lexical entries. The lexicon structure and format of the lexical entries are described in Section 4.

2.3 Trying it out

Assuming you have built the binary `trmorph.fst` using foma, you can simply start foma and use xfst commands implemented in foma to analyze and generate the words. Here is an example session:

```
$ foma
foma[0]: regex @"trmorph.fst";
2.1 MB. 62236 states, 135237 arcs, Cyclic.
foma[1]: up okudum
oku<V><past><1s>
foma[1]: down oku<V><past><2s>
okudun
```

The first line is typed at the shell prompt to start foma. The second line reads the FST specified in `trmorph.fst` into the foma environment. The fourth line asks for the analysis of the verb *okudum* 'I read-PAST', and the fifth line is the output of the analysis. The sixth line asks for the generation of the analysis string produced earlier, modifying the agreement marker to second person singular agreement.

Note that part of the output is removed for readability. We should also note that this example presents one of the rare cases where the analysis is unambiguous. Turkish morphological analysis is an ambiguous process, and TRmorph does not try to avoid it during the analysis. Section 5) lists some of the potential sources of ambiguity. For most purposes, the output of the morphological analyzer needs to be disambiguated. A simple morphological disambiguator is included in TRmorph distribution (see Section 7).

Once you are convinced that the output may be useful for your purposes, you will probably want to use it for analyzing large amount of text. For batch analysis tasks

`flookup` utility distributed with foma is a better fit.

To use the analyzer with the HFST, you need to compile the source automaton with HFST tools.

3 The tagset

The description of the morphology in TRmorph mostly follows Göksel and Kerslake (2005). However, there are some divergences, and tags used in TRmorph analyses does not necessarily match with any of the tags used in any grammar book. This section describes the tags used in the current version of TRmorph. The aim of this section is to help users understand the output of the system. Occasional discussion of the morphological process is included, but this section documents neither the morphology of the language nor the way it is implemented in TRmorph. Our focus in this section is to describe the tags one finds in the analysis strings produced by the analyzer (or tags one needs to use for generation). The index at the end of the document also allows easy access to the points where a particular tag is defined or mentioned in this document.

A clarification of the notation for the surface forms is in order before starting the documentation of the tagset and related suffixes. Suffixes in Turkish often contain under-specified vowels and consonants that are resolved according to morphophonological rules, like vowel harmony. These vowels and consonants are indicated with capital letters listed below.

A is realized as either 'a' or 'e'.

I is realized as either 'ı', 'i', 'u' or 'ü'.

D is realized as either 'd' or 't'.

P is realized as either 'p' or 'b'.

K is realized as either 'k', 'ğ' or 'y'.

C is realized as either 'c' or 'ç'.

A letter in parentheses indicate a buffer consonant or vowel, that may be dropped in certain contexts.

3.1 General structure of the analysis strings

Before describing individual morphological tags used in analysis strings, this section briefly describes the general structure of the analysis strings produced (or accepted in the generation model) by TRmorph. In this document we use the term *morphological tag* for symbols such as ⟨V⟩, or ⟨past⟩. The term *morphological analysis* (or *analysis*) is used for a root word followed by a sequence of morphological tags. In the example presented in Section 2.3, the analysis oku⟨V⟩⟨past⟩⟨1s⟩ (for the word *okudum* 'I read-PAST') consists of the root word *oku* 'read' and morphological tags ⟨V⟩, ⟨past⟩ and ⟨1s⟩ that correspond to part-of-speech category of the root (verb), past tense marker and first person singular subject–predicate agreement marker. The inflections that are default for a word category, such as the fact that the verb above positive (or is not negated), are not indicated in the analyses.

An interesting aspect of Turkish morphology is that the words cannot just be analyzed as belonging to a syntactic category and having a set of inflections based on a category. Due to productive derivational process, an inflected word may change its part of speech and may also get further inflections. Example (1) demonstrates this process with the analysis of of the word *evdekilerinki* the *ones that belong to the ones in the house*, as in 'the *book* that belong to the *people* in the house'.

(1) ev⟨N⟩⟨loc⟩⟨ki⟩⟨Adj⟩⟨0⟩⟨N⟩⟨pl⟩⟨gen⟩⟨ki⟩⟨Adj⟩

The example analysis in (1) can be broken down into the following parts.

1. The initial noun *ev* with the locative maker.

2. Addition of the suffix -ki derives an adjective.

3. The adjective becomes a (pro)nominal with a zero derivation, which is inflected for plural and genitive case.

4. Yet another -ki is suffixed, and the word becomes an adjective again.[3]

The example in (1) is also interesting because of the fact the suffix -ki may result indefinitely long words (see Section 3.5).

---

[3] More likely reading of this example includes another zero derivation causing final POS to be again noun.

In Turkish NLP literature, this process is reflected by so-called *inflectional groups* (IGs) that, for example, can participate in dependency relations. Each step in the above description describes a different inflectional group. The analysis strings produced (or accepted in the generation model) by TRmorph follow the idea of inflectional groups with a slight difference than the examples in the literature. TRmorph makes a distinction between the derivational marker that leads to the POS tag of a IG from the inflectional features of the IG, and the derivational marker is always precedes the POS tag. For example, the second inflectional group in (1) is ⟨ki⟩⟨Adj⟩, indicating the adjective derived by ⟨ki⟩ has not (non-default) inflections.

By default, TRmorph does not mark IG boundaries explicitly. However, one can easily trace the IG changes following the POS tags. All POS tag names start with a capital letter, while other tags always start with a lowercase letter or number. The tag immediately before a new POS tag is always the derivational marker that lead to the new POS tag. If the derivation does not have a corresponding surface affix, a zero-derivation tag ⟨0⟩ is inserted before the POS tag.

3.2 Part-of-speech tags

All part-of speech tags used in TRmorph are listed in Table 1. Most POS tags are self explanatory, and does not require much explanation. The following part of speech tags are somewhat unusual and deserves some explanation.

⟨Exist⟩ is used for two words *var* 'existent/present' and *yok* 'non-existent/absent', where the latter is marked as ⟨Exist:neg⟩, indicating that it is the negative form (see Section 3.3, for the details of this notation). These words behave mostly like nouns in their predicate function (with zero copula), but marking them simply as nouns does blur their function.

⟨Not⟩ is used for *değil* 'not' only. Like *var* and *yok*, *değil* also behaves like nominal predicates. But again, marking it as noun or verb hides the fact that it has a special function.

⟨Q⟩ is used for the question particle -mI. The question particle is written separately from the predicate it modifies. However, the preferred analysis of question

Table 1: The list of part of speech tags in TRmorph.

Tag	Description
⟨Alpha⟩	Symbols of the alphabet
⟨Adj⟩	Adjective
⟨Adv⟩	Adverb
⟨Cnj⟩	Conjunction
⟨Det⟩	Determiner
⟨Exist⟩	The words *var* and *yok*
⟨Ij⟩	Interjection
⟨N⟩	Noun
⟨Not⟩	The word *değil*
⟨Num⟩	Number
⟨Onom⟩	Onomatopoeia
⟨Postp⟩	Postposition
⟨Prn⟩	Pronoun
⟨Punc⟩	Punctuation
⟨Q⟩	Question particle *mI*
⟨V⟩	Verb

particle in TRmorph is together with the predicate. This ensures that it follows the correct form of the predicate it is attached to, and vowel harmony is applied correctly. However, since we do not assume that the input is tokenized with this assumption, this form make sure that the input is analyzed with the cost of precision. The question particle is discussed further in Section 3.15.

3.3 Subcategorization of lexemes

Besides the major major POS tags or word classes discussed above, TRmorph makes use of a set of subcategory tags to mark features that are part of a lexeme. Typically the subcategorization is applied to a root form in the lexicon, but some morphemes and POS tags after a derivation may also receive a subcategory tags. Subcategories defined here are features of a morpheme that do not have a surface realization. Representing these features using a different notation allows one to make this distinction, and the surface–analysis mapping becomes (almost) one-to-one. If a representation where all tags have a uniform notation is desired, the analyzer source can be modified

accordingly, or easier, a simple regular expression based converter can be used.

The subcategories generally mark semantic differences, but they may also result in morphosyntactic differences. Lexical subcategorization in TRmorph output is marked using the syntax ⟨Cat:subcat$_1$:subcat$_2$:…⟩, where 'Cat' is a major category and 'subcat$_1$', 'subcat$_2$' and so on are sub categories. The order of subcategory tags are not important (although they are produced in a consistent order). A typical example of a subcategory is *proper* nouns, which are tagged as ⟨N:prop⟩.

The following lists subcategories used in TRmorph for all word classes that may be specified together with a subcategory.

Nouns Besides the tag ⟨N:prop⟩ marking proper names, abbreviated nouns are marked with the tag ⟨N:abbr⟩. For an abbreviated proper name, the tag is ⟨N:prop:abbr⟩.

Conjunctions are subcategorized as *coordinating*, *adverbial* or *subordinating* conjunctions, marked using tags ⟨Cnj:coo⟩, ⟨Cnj:adv⟩, ⟨Cnj:sub⟩ respectively.

The last one of these categories, ⟨Cnj:sub⟩, include only a limited set of conjunctions which come first in a subordinate clause. These words currently are *ki*, *eğer* and *şayet* (all borrowings from Persian). The other subordinating particles/words occur at the end of subordinate clauses, and they are marked as postpositions (⟨Postp⟩) described below. Furthermore, most of the subordination in Turkish is done through suffixation which is described in Section 3.16.

Pronouns Pronouns are further categorized as *personal*, *demonstrative* and *locative* pronouns, marked using ⟨Prn:pers⟩, ⟨Prn:dem⟩, ⟨Prn:locp⟩ respectively. Furthermore, the pronouns that form questions, like *kim* 'who', and *ne* 'what', are marked as ⟨Prn:qst⟩. Subcategory markers for both aspects can be present. For example *kim* 'who' would be marked as ⟨Prn:pers:qst⟩.

Besides the above subcategories, personal pronouns get person-number agreement markers. These markers can be useful in subject-predicate agreement as well as in other constructions (such

as genitive-possessive construction involving pronouns). However, the agreement in Turkish is far from trivially determined (see Göksel and Kerslake (2005, pp.116–122)). The markers ⟨Prn:pers:1s⟩, ⟨Prn:pers:2s⟩, ⟨Prn:pers:3s⟩, ⟨Prn:pers:1p⟩, ⟨Prn:pers:2p⟩ and ⟨Prn:pers:3p⟩ are tags used for the personal pronouns with person-number agreement. The agreement markers are further discussed in Section 3.13.

The reflexive pronoun *kendi* and its different person forms are marked as ⟨Prn:refl⟩. Like other personal pronouns, reflexive pronouns are also marked with a person agreement marker.

Subcategorization of pronouns, particularly as personal pronouns, are sometimes not a clear decision. Subcategories of some pronouns are left unspecified even though they are often used as personal pronouns, and some pronoun marked as personal pronouns may refer to entities other than people.

Determiners are marked for definiteness. *Definite* determiners are marked ⟨Det:def⟩ and *indefinite* determiners are marked ⟨Det:indef⟩. The question words that fill the same syntactic slot as determiners *ne kadar* 'how much' and *hangi* 'which' are tagged with ⟨Det:qst⟩.

Further subcategorization of determiners (for example quantifiers) can be implemented in the future.

Postpositions are always subcategorized in two dimensions. First subcategory is the syntactic category (POS) of the resulting postpositional phrase, either an *adjectival* or *adverbial* phrase, marked as ⟨Postp:adj⟩ and ⟨Postp:adv⟩ respectively. Note that unlike other POS tags, these category markers start with a lowercase letter.

Postpositions choose their noun phrase complements. Besides the category of the resulting phrase, postpositions also include a tag specifying the requirement for the complement noun phrase. The tag marking required complement type is formed by a concise description of the requirement followed by the capital letter 'C'. The postpositions that require the complement to be in *ablative*, *accusative*, *dative*, *genitive* and *instrumental*

cases are marked ⟨Postp:ablC⟩, ⟨Postp:accC⟩, ⟨Postp:datC⟩, ⟨Postp:genC⟩, and ⟨Postp:insC⟩ respectively. The postpositions that require the noun phrase to be suffixed with either -lI or -sIz are marked with ⟨Postp:liC⟩.[4] Postpositions that require non-case marked complement are tagged ⟨Postp:nomC⟩. Finally, postpositions that require numeric expressions as their complements are marked with ⟨Postp:numC⟩. For some the postpositions that take more than one type of noun complements, TRmorph produces only the (presumably) most common option. For example, the postpositions that are marked as ⟨nomC⟩ also take genitive marked pronouns as complements. Similarly, postpositions *önce* and *sonra* that normally take ablative complements, can also take bare (non-case-marked) numbers or time expressions.

Numbers are tagged as ⟨Num:ara⟩ for Arabic numerals, and ⟨Num:rom⟩ for Roman numerals. Numbers that are spelled out are not marked with a subcategory marker (but still marked as ⟨Num⟩). Besides numbers, the question word *kaç* 'how many' is also tagged as a number with a sub tag specifying that it is a question word, resulting in ⟨Num:qst⟩.

A rudimentary date/time recongnition is also supported. Time expressions in 24-hour format with optional 'seconds' field, i.e, hh:mm(:ss) (also with separator '.'), and date formats of DD/MM/YYYY (also with separators '.' and '-' and YYYY-MM-DD are recognized as dates and times, and tagged as ⟨Num:time⟩ and ⟨Num:date⟩, respectively.

Verbs are currently not subcategorized in TRmorph.

Subcategorizing verbs as *transitive* and *intransitive*, or marking all types (cases) of noun phrase complements a verb can take is planned and some early steps are underway as of this writing (July 2013).

Adverbs are not currently subcategorized, except a few adverbial question words for which the tag ⟨Adv:qst⟩ is used.

Exist The tag ⟨Exist⟩ exists only for two words *var* 'existent/present' and *yok* 'non-existent/absent'. Since

---

[4]These suffixes are typically considered derivational suffixes, however their use resemble case markers.

yok is the negative of *var*, it is tagged as negative: $\langle$`Exist:neg`$\rangle$.

Some verbs, nouns, adjectives, adverbs and conjunctions are formed by more than one written words. Some of these are adjacent words, like the adverb *apar topar* 'hurriedly', but some may be split like the conjunction *ya*, as in *ya evdedir ya iş yerinde* 's/he is either at home or the office'. Furthermore, some of individual 'words' in such constructions cannot be used by themselves, like *topar* above. If the non-split multi-word expressions are input to the analyzer together, they are analyzed like other words of the same class. However, if they are input word-by-word, a sub tag $\langle$`:partial`$\rangle$ is added to the main POS tag. For example *apar* and *topar* are tagged as $\langle$`Adv:partial`$\rangle$ and *ya* is tagged as $\langle$`Cnj:partial`$\rangle$ (more precisely $\langle$`Cnj:coo:partial`$\rangle$). Currently, the tags $\langle$`N:partial`$\rangle$, $\langle$`Adj:partial`$\rangle$ and $\langle$`V:partial`$\rangle$ are used for parts of nouns, adjectives and verbs respectively.

3.4 Nominal morphology and noun inflections

Nouns, pronouns, adjectives and adverbs in Turkish form the larger class of nominals. Most adjectives, and some adverbs can function as nouns (or pronouns). For example, *mavi* 'blue' may have a noun reading 'the blue one'. Similarly, some adverbs like *şimdi* 'now' may take nominal inflections *şimdilerde* 'now-PL-LOC = (literally) in current times'. In TRmorph this is handled by allowing any adjective or adverb to become an noun with a *zero derivation*.[5] A zero derivation is always marked with the tag $\langle$`0`$\rangle$ followed by the new POS tag, in this case $\langle$`N`$\rangle$.

Nouns can be suffixed with the plural suffix, one of the possessive suffixes and one of the case suffixes. All of these inflections are optional. When not marked with any of these suffixes, the default is singular, no possessive marking, an no case marking (or nominal), respectively. When these suffixes co-occur, they have to occur in the order listed, shown in Figure 1. The full list of noun inflections are presented in Table 2.

If there is a plural marker, analysis string after the $\langle$`N`$\rangle$ will include the tag $\langle$`pl`$\rangle$. TRmorph does not mark for

---

[5]This certainly generates incorrect analyses for a large number of adverbs which do not 'nominalize'.

Figure 1: Automata depicting noun inflections. The edge CASE1 represents the locative and ablative suffixes, CASE2 represents all other case-like suffixes. The reason for the differentiation is due to the fact that the state CASE1 can be followed by the suffix -ki.

singular. If a noun is not marked for plural, it is assumed to be singular.

The first five suffixes in the lower part of Table 2 are commonly recognized cases in Turkish. The *instrumental/commutative* marker also behaves like case suffixes. There are two more suffixes, namely -ll and -slz that can occupy the same slot, which are marked with tags $\langle$`li`$\rangle$ and $\langle$`siz`$\rangle$ respectively.

Possessive markers follow either the nominal stem, or the plural marker. The basic function of the possessive markers are to mark a noun for possession. That is a noun belonging to some entity, e.g., *evi-m* 'my house' or *ev-i* 'his/her house'. Besides marking for possession, these suffixes, particularly the third person possessive suffix, have a number of other functions. The rest of this section explains some of these usage patterns, and how TRmorph represents them.

TRmorph normally does not allow adjectivals (adjectives, determiners and numbers) to take any of the possessive suffixes directly. However an adjectival suffixed one the possessive suffixes may function as a pronoun. Examples include, *üç-ümüz* 'three of us', *bazı(lar)-ınız* 'some of you' and *eski-si* 'the old one (of them)'. Note that this usage is different than possessively marked adjective with the noun interpretation, e.g., not 'the three that belongs to us' but 'three of us'. In this use, possessive markers are treated like a deriva-

Table 2: Noun inflections.

Function		surface	tag
Plural		-lAr	⟨pl⟩
Possessive	First person singular	-(I)m	⟨p1s⟩
	Second person singular	-(I)n	⟨p2s⟩
	Third person singular	-(s)I	⟨p3s⟩
	First person plural	-(I)mIz	⟨p1p⟩
	Second person plural	-(I)nIz	⟨p2p⟩
	Third person plural	-lArI	⟨p3p⟩
Case	Accusative	-(y)I	⟨acc⟩
	Dative	-(y)A	⟨dat⟩
	Ablative	-DAn	⟨abl⟩
	Locative	-DA	⟨loc⟩
	Genitive	-(n)In	⟨gen⟩
	Instrumental/commutative	-(y)lA	⟨ins⟩

tional suffix. The examples above would be analyzed as üç⟨Num⟩⟨p1p⟩⟨Prn⟩, bazı⟨Det:indef⟩⟨p1p⟩⟨Prn⟩ and eski⟨Adj⟩⟨p3s⟩⟨Prn⟩, respectively.

A similar usage is observed with verbal nouns and participles (see Section 3.16). In these cases the possessive marker marks the subject of the verb. For example, in participle use of *oku-yacağ-ım* '(the book) that I will read', the possessive suffix marks who does the reading, and not a possession relation in the usual sense. Currently trmorph analyzes this word as oku⟨V⟩⟨part:fut⟩⟨Adj⟩⟨p1s⟩.

The -(s)I suffix, listed as ⟨p3s⟩ in Table 2, is highly ambiguous. One of its many functions that may be confused with the possessive suffix is forming noun compounds. In earlier versions of TRmorph, this function of -(s)I was always marked with the tag ⟨ncomp⟩. This marker can be useful for marking noun compounds like *at araba*sı 'horse carriage'.[6] In this use, this tag always causes ambiguities. Besides the fact that a noun suffixed with -(s)I can either be marked for possession or as the head of a noun compound, since one of the two -(s)I suffixes following each

other is deleted from the surface form, it can also be both (a noun compound marked for possession, *at araba*sı '*his* horse carriage'). In case any or the other possessive markers are used with a noun compound, the suffix -(s)I is again deleted (e.g., *at araba*nız '*your* horse carriage'). In summary, marking heads of nominal compounds are not straightforward during the analysis. As a result this marker is a compile time option in the current version (disabled by default). If not enabled, one should note that the tag ⟨p3s⟩ may indicate a compound head with or without third person singular possessive marking (see also the discussion of ambiguity regarding ⟨p3s⟩ and ⟨p3p⟩ tags below).

Another issue with the -(s)I suffix is that a noun marked with -(s)I may also indicate a third person plural possessor, e.g., *onların araba*sı '*their* car'. In general, if there is an overt possessor, the preferred third person plural marker is -(s)I, rather than -lArI. TRmorph marks -(s)I both as ⟨p3s⟩ and ⟨p3p⟩.

The case (or case-like) suffixes change the role of the noun (or the noun phrase headed by the noun) in the sentence. For example a locative marked noun phrase may function as an adverb (*saat dokuz*da *görüşürüz*) or an adjective (*yedi yaşın*da *çocuk*). However, following the common practice in the literature we do not attempt to mark possible POS changes after case-like markers.

3.5 The suffix -ki

The suffix -ki, tagged as ⟨ki⟩, attaches to locative or genitive marked nouns. The suffix may also attach to nouns expressing (a unit of) time, e.g., *ay-ki* 'monthki'.[7] The resulting word functions as an adjective or a pronoun. In both cases, TRmorph marks the transition to an adjective. For example, *evdeki* is analyzed as 'ev⟨N⟩⟨loc⟩⟨ki⟩⟨Adj⟩'. Since all adjectives are allowed to become a noun through a zero derivation, the pronoun reading is intended to be represented by this change. For example, the intended analysis for *evdeki kitap* 'the book in the house' is 'ev⟨N⟩⟨loc⟩⟨ki⟩⟨Adj⟩', while analysis for *evdeki uyuyor* 'the one/person in the house is sleeping' appends ⟨0⟩⟨N⟩ at the end of the analysis string.

The (pro)noun formed by -ki can further be suffixed

---

[6]Even though one can assume that this use is somewhat related to possession, it is not strictly possessive marking (the horse does not own the carriage). Furthermore, since a -(s)I after another one is deleted on the surface, a single -(s)I suffix may also indicate a nominal compound in possessive form (e.g., 'someone's horse carriage').

[7]In this use, the suffix affects a larger 'time phrase', like *bu yılki üretim* 'this-year's production'.

with other nominal suffixes. Although the number of iterations using -ki rarely exceed two in practice, there is no principled limit. As a result, length of a Turkish word is in-principle unbounded.

3.6 Tags related to nominal predicates

Any nominal in Turkish may become a predicate with one of the copular suffixes -(y)DI, -(y)mIş, -(y)sA or -(y). These suffixes correspond to *past*, *evidential*, *conditional*, and *present* predicates involving the copula 'be'. The copular markers has to precede one of the verbal person agreement markers. For example *öğrenciydik* 'we were students', *öğrenciymişler* 'they were [evidentially] students', *öğrenciysen* 'if you are/were a student', *öğrenciyim* 'i'm a student'. Since the third person singular agreement suffix is null on the surface and the buffer -(y)- does not surface in this case, any nominal without additional copular or person suffixes serve as a nominal predicate with present copula and third person singular agreement. Additionally, since a predicate with third person singular agreement also agrees with a third person plural subject, we additionally mark such a noun as having present copula and third person plural agreement (for example, *babam öğretmen, annem ve ablam doktor* 'my father is a teacher, my mother and older sister are doctors').

TRmorph handles this process by allowing any noun and adjective to first became a verb with a zero derivation, and then marking it with the appropriate copula and the person agreement marker. The tags for copula are ⟨cpl:pres⟩, ⟨cpl:past⟩, ⟨cpl:evid⟩ and ⟨cpl:cond⟩ for *present*, *past*, *evidential* and *conditional* copula respectively. Last three tags are also possible after a verb with a tense/aspect/modality suffix, and is discussed further in Section 3.14. Example analyses for the examples discussed above would be as follows:

öğrenciydik	⟨N⟩⟨0⟩⟨V⟩⟨cpl:past⟩⟨1p⟩
öğrenciymişler	⟨N⟩⟨0⟩⟨V⟩⟨cpl:evid⟩⟨3p⟩
öğrenciysen	⟨N⟩⟨0⟩⟨V⟩⟨cpl:cond⟩⟨2s⟩
öğrenciyim	⟨N⟩⟨0⟩⟨V⟩⟨cpl:pres⟩⟨1s⟩
öğretmen	⟨N⟩⟨0⟩⟨V⟩⟨cpl:pres⟩⟨3s⟩
doktor	⟨N⟩⟨0⟩⟨V⟩⟨cpl:pres⟩⟨3p⟩

Besides copular suffixes, the suffix -(y)ken (making adverbials from verbs, discussed in Section 3.16) may occupy the same slot as the copular suffixes, although its use is more restricted.

The nominal predicate with a copula and person agreement may be followed by the marker Göksel and Kerslake (2005) call 'generalizing modality marker', the suffix -DIr. It is particularly common with ⟨3s⟩ as it disambiguates between the noun and the predicate reading. The tag for this marker in TRmorph is ⟨dir⟩.

3.7 Number inflections

The suffix -(ş)Ar, tagged ⟨dist⟩, attached to numbers form *distributive* numerals. Besides the numbers (written as numerals or spelled out), question word *kaç* 'how many' may also get this suffix, and tagged with ⟨dist⟩.

The ordinal numerals are formed using the suffix −(I)ncI, and tagged as ⟨ord⟩. Ordinals are also specified by a 'dot' after Arabic or Roman numerals. TRmorph currently does not handle this notation.

Percent sign before a numeral is treated like a prefix, and tagged as ⟨perc⟩.

3.8 Apostrophe behavior

In written text an apostrophe is required after proper nouns and numbers (official rules are more complicated). However, the real-world use is rather relaxed, and people tend to omit the apostrophe.

Another difficult case for apostrophe is after the compound proper nouns, like *Türkiye Büyük Millet Meclisi* 'Grand National Assembly of Turkey', *Ağrı Dağı* 'Mount Ararat' or *Öfkeli Şirin* 'Grouchy Smurf'. Unless tokenized together, the analyzer cannot know that these words are part of a proper noun, and parts of these compounds will be tagged as if they are single words. If the last noun in a compound is part of a proper noun, an apostrophe is required if further suffixes follow the last noun. TRmorph allows bare nouns, nouns with an ⟨ncomp⟩ tag or when ⟨ncomp⟩ is not enabled, nouns with a ⟨p3s⟩ tag to have an optional apostrophe before other suffixes. This behavior can be disabled during compile time in `options.h`.

3.9 Verbal voice suffixes

Turkish verbs can be suffixed with one or more of the voice suffixes *reflexive*, *reciprocal*, *causative* and *passive*. The tags used for these functions are ⟨rfl⟩, ⟨rcp⟩, ⟨caus⟩

Table 3: Suffixes that make compound verbs.

Suffix	Tag	Expresses
-(y)Abil	⟨abil⟩	ability
-(y)Iver	⟨iver⟩	immediacy
-(y)Agel	⟨agel⟩	habitual/long term
-(y)Adur	⟨adur⟩	repetition/continuity
-(y)Ayaz	⟨ayaz⟩	almost
-(y)Akal	⟨akal⟩	stop/freeze in action
-(y)Agör	⟨agor⟩	somewhat like ⟨iver⟩

Table 4: Tense/aspect/modality markers. The usage of suffix -(y)A to express conditional aspect is informal, and rather restricted. Aorist suffix is highly irregular. The choice of -Ar and -Ir depends on the stem. The -z form occurs only after negative marker, and it is not realized on the surface if it precedes first person agreement suffixes.

Tag	Suffix	Description
⟨evid⟩	-mIş	evidential past (perfective)
⟨fut⟩	-(y)AcAk	future
⟨obl⟩	-mAlI	obligative
⟨impf⟩	-mAktA	imperfective
⟨cont⟩	-(I)yor	imperfective
⟨past⟩	-DI	past (perfective)
⟨cond⟩	-sA,-(y)A	conditional
⟨opt⟩	-(y)A	optative
⟨imp⟩	-	imperative
⟨aor⟩	-Ar,-Ir,-z,-	aorist

and ⟨pass⟩, respectively. The first two are rather unproductive while causative and passive forms are productive. Furthermore, causative suffix can be used repetitively.[8] With some verbs, use of double causative suffix yields the same semantics as a single causative suffix. TRmorph does not treat these cases separately. If surface string has double causative suffixes, the analysis will include two ⟨caus⟩ tags, regardless of its semantics.

Despite the fact that most grammar books list voice suffixes under inflectional morphology, TRmorph treats them as derivations, i.e., a ⟨V⟩ tag follows the voice related tags.

3.10 Compound verbs

A verbal stem (possibly including voice suffixes) may be followed by a set of suffixes listed in Table 3 to form compound verbs. These suffixes are related to some standalone verbs.

The first three suffixes in this Table 3 are relatively productive, the others are rare or their use are mostly lexicalized. Although not frequent in use, more than one these suffixes may attach to the same stem, for example *çıkıverebilir* 'he/she/it may possibly come out/show up' analyzed as 'çık⟨V⟩⟨iver⟩⟨V⟩⟨abil⟩⟨V⟩⟨aor⟩⟨3s⟩'.

The form of ⟨abil⟩ in a negative verb is -(y)A, and unlike the rest of the suffixes listed in Table 3 it follows the negative marker.

Like the voice suffixes, we treat these suffixes as derivations, starting a new verbal inflectional group.

3.11 The negative marker

Negation of a verbal predicate is indicated with the suffix -mA, and marked simply as ⟨neg⟩. Nominal predicates do not get this suffix, instead the particle *değil* is used.

3.12 Tense/aspect/modality markers

A verb with a set of suffixes described above either becomes a finite verb by taking one of the tense, aspect and modality (TAM) markers followed by a person-number agreement suffix, or it can be subject to subordination and becomes nominalized.

The list of TAM suffixes, the corresponding tags and brief descriptions are given in Table 4.

3.13 Person and number agreement

After TAM markers a finite verb requires one of the person and number agreement markers. For any finite predicate an agreement marker is compulsory.

Morphological tags for the agreement markers are composed of two characters, e.g., ⟨1s⟩ or ⟨3p⟩. The first character is a number indicating the person (**1**st, **2**nd or **3**rd),

---

[8]Again, although this is limited in practice, there is no principled limit on the number of causative suffixes that one can string one after another.

9

Table 5: Person-number agreement markers. The suffixes listed in the column marked 'TAM1' follow the TAM markers ⟨evid⟩, ⟨fut⟩, ⟨obl⟩, ⟨impf⟩ and ⟨cont⟩ as well as the evidential copula ⟨cpl:evid⟩ and nominal predicates. The same set of suffixes also follow positive verbs with ⟨aor⟩ without a negative marker. The suffixes on the column marked 'TAM2' are used after ⟨past⟩ and ⟨cond⟩ as well as the corresponding copular markers ⟨cpl:past⟩ and ⟨cpl:cond⟩.

Tag	TAM1	TAM2	optative	imperative
⟨1s⟩	-(y)Im	-m	-(y)Im	*
⟨2s⟩	-sIn	-n	-sIn	-
⟨3s⟩	-	-	-	-sIn
⟨1p⟩	-(y)Iz	-K	-lIm	*
⟨2p⟩	-sInIz	-nIz	-sInIz	-(y)In,-(y)InIz
⟨3p⟩	-lAr	-lAr	-lAr	-sInlAr,-

and second one indicates the number (**s**ingular or **p**lural). The surface form of the person-number agreement markers change depending on the suffixes they follow. Table 5 lists the person agreement markers and their surface forms according the TAM of the verb they attach to. Note that the third person singular marker is null on the surface after most TAM markers. Furthermore, since a predicate with third person singular marker will also agree with third person plural subject, all forms that are marked with a ⟨3s⟩ tag will also be marked with a ⟨3p⟩ tag. A predicate with plural agreement marker may also indicate or agree with a singular subject in polite/formal speech. However, TRmorph does not mark predicates with plural agreement markers with the ⟨3s⟩ tag.

Although an agreement marker is required for a predicate, TRmorph accepts a predicate with a TAM marker but without an agreement marker.[9] There are a few cases where the agreement marker can be attached to somewhere else. In a relatively common case, the agreement marker may be on the question particle (see Section 3.15), as in ⟨2s⟩ suffix -sIn in *(partiye) gidecek-mi**sın**?* 'will you go (to the party)?' (dash '-' indicates

the expected position of the suffix in bold case). Similarly, a verb may lack an agreement marker if it is part of a compound verb formed by free auxiliaries *ol* and *bulun*, for example *gitmiş- bulun**duk*** 'we happened to go', or *gidecek- olursa**nız** (üzülürüm)* 'if you were to go (I'd be upset')'.[10] Another case for predicates without agreement markers is 'delayed suffixation' of coordinated predicates, where only last predicate may have the agreement marker. For example *kahvaltı eder-, oltalarımızı alır-, balığa gider**iz*** 'We have breakfast, take our hooks, and go fishing'. Note that the predicates without agreement markers will also receive ⟨3s⟩ and ⟨3s⟩ analyses, since a predicate with a null agreement suffix agrees with third person singular and plural subjects. However, in all the examples above, tagging the main the predicates with third person markers would be incorrect.

One final note on agreement markers is that the subject–predicate agreement in Turkish is more complicated than simply matching agreement markers on subject noun phrase and the predicate (see Göksel and Kerslake 2005, chapter 5). TRmorph marks nominals for number (if it is plural), but not for person, except in a few lexically specified cases, e.g., personal pronouns. It would be easy to modify TRmorph to mark any noun that can become a subject (typically no case marked nouns, but also genitive in some cases) with third-person agreement. However, utility of such a marker is not clear.

3.14 Copular markers and -**DIr**

The copular suffixes discussed in Section 3.6 can also be attached to a verb after a TAM marker, typically forming complex tenses. These suffixes are -(y)DI, -(y)mIş and -(y)sA, tagged as ⟨cpl:past⟩, ⟨cpl:evid⟩ and ⟨cpl:cond⟩, respectively.

The conditional copula -(y)sA can co-occur with other copular markers. When there is a copular suffix, person-number agreement suffixes normally attach after the first copula. However the third person plural suffix may be after the TAM marker or second copular suffix as well.

Similar to the nominal predicates with a copula, copular suffixes may be followed the 'generalizing modality marker' -DIr tagged as ⟨dir⟩.

3.15 The question particle

Question particle -mI, tagged as ⟨Q⟩, is normally written separately. However, it has an intimate relationship with the verb or the nominal predicate it attaches to. First, a few exceptions aside, it is attached to a tensed verb without a person agreement. In this case, the person agreement and the suffixes that may follow must be attached to the question particle. In this particular case, the verb will often be analyzed wrongly as having the agreement marker ⟨3s⟩ or ⟨3p⟩, since a predicate with null person agreement suffix may agree with third person singular or plural subjects. Second, the question particle follows the vowel harmony rules, and the underspecified vowel on -mI is realized based on the last vowel of the verb. As a result the question particle can only be analyzed (and generated) with precision only together with the word it is attached to.

If tokenized together with the predicate, TRmorph will swallow the space in between the predicate and the -mI and analyze it altogether. In this case the lowercase tag ⟨q⟩ is used. Furthermore, it is a common spelling mistakes to write the question particle together with the related word. TRmorph can be instructed to to accept this common mistake during compile time, in which case the tag will again be ⟨q⟩.

3.16 Subordination

A set of suffixes attached to an 'untensed' verb, a verb without any TAM markers, result in the phrase headed by the verb to become a subordinate clause. TRmorph follows the description in Göksel and Kerslake (2005), and makes the distinction between three different forms of subordination. First, a set of suffixes produce *verbal nouns* from a non-finite verb. The resulting words function as the head of the noun phrases, and with some limitation they can receive all nominal inflections. The second group forms *participles*, which form relative clauses. Participles can also take nominal inflections with few restrictions. The last group, *converbs*, form adverbials and they are more restricted in terms of the morphemes attached

Table 6: Subordinating suffixes and tags used for subordinating suffixes.

Tag	Suffix
⟨vn:inf⟩	-mA
⟨vn:inf⟩	-mAK
⟨vn:yis⟩	-(y)Iş
⟨vn:past⟩	-DIk
⟨vn:fut⟩	-(y)AcAk
⟨vn:res⟩	-(y)An
⟨part:past⟩	-DIk
⟨part:fut⟩	-(y)AcAk
⟨part:pres⟩	-(y)An
⟨cv:ip⟩	-(y)Ip
⟨cv:meksizin⟩	-mAksIzIn
⟨cv:ince⟩	-(y)IncA
⟨cv:erek⟩	-(y)ArAk
⟨cv:eli⟩	-(y)AlI
⟨cv:dikce⟩	-DIkCA
⟨cv:esiye⟩	-(y)AsIyA
⟨cv:den⟩	-dAn
⟨cv:den⟩	-zdAn
⟨cv:cesine⟩	-CAsInA
⟨cv:ya⟩	-(y)A
⟨cv:ken⟩	-(y)ken

to them. The suffixes that form forms different types of subordinating suffixes overlap significantly. As a result, producing ambiguous analyses.

TRmorph uses the tag structure ⟨type:subtype⟩ for marking subordinating suffixes. The first part, type, is one of vn, part and cv for verbal nouns, participles and converbs, respectively. The second, subtype, part indicate a further distinction of the function of the suffix, a relevant linguistic abbreviation, but sometimes a version of the surface form of the suffix. The tags used for all three types of subordinating suffixes are listed in Table 6.

Since verbal nouns, participles and converbs derive nominal, adjectival and adverbial phrases, respectively, POS tags, ⟨N⟩, ⟨Adj⟩ and ⟨Adv⟩, follow these tags.

Some of the suffixes have multiple functions and may derive more than one type of subordinate clauses. Fur-

thermore, TRmorph will produce some spurious ambiguity because of the fact that any adjective, hence a word suffixed with an participle, is allowed to become a noun with a zero derivation.

The list in Table 6 follows Göksel and Kerslake (2005). The main exception is the suffixes listed by Göksel and Kerslake (2005) as converbial suffixes that require a post-position. Since the postposition in these cases will signal the adverbial function of postpositional phrase, TRmorph does not mark the complement of the postposition as a converb.

Most of these suffixes attach to an untensed verb. Except, the suffix -(y)ken which behaves much like the copular suffixes discussed above. Furthermore, the -(y)A in its subordinating function is typically used together with reduplication, e.g., *koşa koşa* 'run-(y)A run-(y)A = hurriedly', but also occurs in words like *diye*, where it does not need reduplication.[11]

Besides the subordinating suffixes (participles) discussed above, some of the TAM markers, namely ⟨aor⟩ (-Ar/-Ir), ⟨evid⟩ (-mIş) and ⟨fut⟩ (-AcAk), derive adjectives from verbs.[12] TRmorph handles this by analyzing any verb with one of these TAM markers without further suffixes (e.g., agreement markers) as an adjective. For example, the word *görülmüş* in *görülmüş mektup* 'see-PASV-EVID letter = the letter that was seen' is analyzed as '`gör⟨V⟩⟨pass⟩⟨V⟩⟨evid⟩⟨Adj⟩`'.

3.17 Productive derivational morphemes

Almost all the tags and relevant morphological process above are described as part of inflectional morphology in most grammar books. The suffixes described here are the ones that are traditionally considered derivational suffixes. Some of these suffixes, for example -lI and -slz discussed earlier, may attach to word forms that are already inflected by other suffixes. Others normally attach only to the stem and produce another stem.

The derivational suffixes analyzed by TRmorph are listed in Table 7. Of these suffixes, -lA causes a large number of ambiguous analyses since it is part of a

few other suffixes, for example, the plural suffix -lAr. In case of -lAr, whose remainder -r also matches a verbal suffix (aorist marker), any plural noun is analyzed as a verb as well. For example, the word *kanıtlar* is analyzed as `kanıt⟨N⟩⟨pl⟩` 'evidence-PLU', or `kanıtN⟨la⟩⟨V⟩⟨aor⟩⟨3s⟩` '(s)he proves (it)'. Hence, including -lA in the analysis causes an increase in the number of analyses produced by TRmorph. The analysis of the suffix -lA can be disabled during the compile time. In this case, a separate lexicon containing verbs with derivational suffix -lA will be included in the root lexicon. Some of the 'lexicalized' verbal stems that end in -lA are included in the main lexicon, and will result in an analysis with a stem including -lA regardless of this option. Similarly, -lA is always accepted on onomatopoeic roots and interjections.

In some cases the derivational suffixes can only be attached to a subset of the words within the word class. Not all these cases are easy to identify, and TRmorph will analyze combinations of stems and suffixes that are not in the language. Furthermore, TRmorph does not limit the number of derivational suffixes that can be stringed one after another other, even though multiple derivations of this sort is a lot more restricted.

Besides the sources of possible erroneous over-analyses listed above, the derivational morphology specification in TRmorph over-generates in some cases. In particular, any form of the diminutive suffix is allowed to attach to any noun, although most nouns are used only one of the diminutive suffixes. The ambiguity and overgeneration are discussed in Section 5.

4 The lexicon

TRmorph contains a root lexicon which is created extracting root forms from a large web corpora, and checking the possible forms against online dictionaries, and the lexicon of the earlier version which was based on Zemberek (A. A. Akın and M. D. Akın 2007). The result is also checked and corrected manually as part of the development process.

The lexicon files are located under the directory `lexicon` and included (through C preprocessor) as a single root lexicon. The files under `lexicon/` are simply a list of root forms and their *continuation classes*. Continuation classes can be any `LEXICON` declaration in the file `morph.lexc`, but typical continuation classes are the main

---

[11] We also analyze *diye* as a postposition, as its use as subordinator is semantically unlike the others uses of -(y)A.

[12] Although they differ in usage, these forms have the same semantics as the verb-verb compounds formed by the auxiliary auxiliary verb *ol* with present participle suffix (*ol-an*).

word (POS) categories, such as `N`, `Adj` and `V`. The lexical exceptions are specified after the main category information. For example, `V_AorAr` for verbal roots that take the exceptional -Ar form of the aorist suffix. Likewise, `N_comp` is used for lexicalized nominal -sI compounds since when these words are pluralized the plural marker is inserted between the word and the suffix -sI.

The lexical forms are similar to the written forms of the relevant stem. However, a set of special 'multi-character' symbols are used for providing information necessary for morphophonological processing. A large group of these symbols are concerned with 'final stop devoicing' (or voicing depending on your view point). The consonants ç, t, k, p and g at the end of some of the roots are replaced with their voiced counterparts if they precede a suffix that starts with a vowel. These root forms are lexically marked by replacing the consonants above with multi-character symbols ^c, ^t, ^k, ^p and ^g, respectively.

Besides the voicing changes of consonants, some borrowings end with a 'palatalized' consonant that affects vowel harmony process. For example *saat* 'watch/clock' is inflected as *saat-i* 'watch-ACC' instead of *saat-ı* as vowel harmony suggests. These words are indicated by the vowel before such a consonant by a three-letter multi-character symbol. These symbols always start with ^p and a capitalized version of the relevant vowel. For example, the word *saat* is listed as `sa^pAt` in the lexicon.

One last class of similar special symbols are so-called *silent vowels and consonants*. These are particularly useful for abbreviations and numerals, but also some names of foreign origin. The suffixes that follow such words are also subject to morpho-phonological process like vowel harmony. However, this cannot be derived from their written form. For example correct inflected form of *ABD-DAT* 'USA-DAT' is *ABD'ye*, not *ABD'ya*. The way to solve this problem is to insert a silent (front-unrounded) vowel after the abbreviated form. The multi-character symbols ^sBUV ^sBRV ^sFUV ^sFRV ^sVC and ^sUC are used for silent vowels and consonants (see the comments in file `lexicon/abbreviation` for more information).

A somewhat inconsistent notation is used for three morphological processes. First, the multi character symbol `@DEL@` is inserted before a vowel that is deleted if a suffix starting with a vowel follows. Second, the last consonant in some borrowings are duplicated if they follow a suffix that start with a vowel. These root forms are marked by

inserting the multi-character symbol `@DUP@` before the duplicated consonant. And the last symbol `@DELS@` is used in lexical entries of a few borrowed words which delete s in the suffix -sI.[13]

5 Ambiguity and overgeneration

This section discusses the ambiguous analyses in TRmorph, and also touches upon a related but different problem, overgeneration.

The morphological analysis of Turkish text is an inherently ambiguous process. However, the design choices made in a morphological analyzer affects the number of ambiguous analyses produced per word. TRmorph, by design, does not try to reduce the number of ambiguous analyses. In general, TRmorph produces more ambiguous analyses than the others (mainly based on Oflazer (1994)) reported in the literature.

The following is a list of cases where one finds ambiguous morphological analyses in TRmorph. Some of these cases are not specific to TRmorph, and for example, noted by Oflazer and Tür (1997) as well. This list may be useful for the users who may wish to disambiguate the output of the analysis using rule-based methods, or it may also be useful in the process of designing statistical disambiguators.

1. Ambiguous root forms, for example *yüz* can be analyzed as:

   (a) `yüz⟨N⟩` 'face'
   (b) `yüz⟨Num⟩` 'hundred'
   (c) `yüz⟨V⟩⟨imp⟩⟨2s⟩` 'swim'

2. A root form is the same as a shorter root and one or more suffixes, for example *buna* can be analyzed as

   (a) `bu⟨Prn:dem⟩⟨dat⟩` 'this-DAT'
   (b) `buna⟨V⟩⟨imp⟩⟨2s⟩` 'become senile-IMP'
   (c) `bun⟨N⟩⟨dat⟩` 'trouble-DAT'

---

[13] These multi-character symbols are both inconsistent with the others, and they may be confused with 'flag diacritics' at first sight (TRmorph does not use any flag diacritics). This notation in the lexicon may change in the future version of TRmorph.

Note that the root 'bun' is a very rare/regional word, and the imperative verb reading is also very unlikely. However the best option for the analyzer is to produce all these analyses, and let the later stages analysis disambiguate between them.

3. The surface form of a suffix is a combination of two other suffixes. For example, the word *evleri* can be

    (a) ev-leri 'ev$\langle$N$\rangle\langle$p3p$\rangle$ = their house'
    (b) ev-ler-i 'ev$\langle$N$\rangle\langle$pl$\rangle\langle$acc$\rangle$ = houses-ACC'

    Furthermore, the same word can also be analyzed as

    (a) ev$\langle$N$\rangle\langle$pl$\rangle\langle$p3s$\rangle$
    (b) ev$\langle$N$\rangle\langle$pl$\rangle\langle$p3p$\rangle$
    (c) ev$\langle$N$\rangle\langle$ncomp$\rangle\langle$p3p$\rangle$
    (d) ev$\langle$N$\rangle\langle$ncomp$\rangle\langle$pl$\rangle$
    (e) ev$\langle$N$\rangle\langle$ncomp$\rangle\langle$pl$\rangle\langle$p3p$\rangle$
    (f) ev$\langle$N$\rangle\langle$ncomp$\rangle\langle$pl$\rangle\langle$p3s$\rangle$
    (g) ev$\langle$N$\rangle\langle$ncomp$\rangle\langle$pl$\rangle\langle$p3p$\rangle$

    The reason for these analyses has to do with the sources of ambiguity explained in items 6 and 8.

4. An analysis with multiple morphemes is also a (derived) lexicalized form. For example the word *konuşma* can be analyzed as

    (a) konuşma$\langle$N$\rangle$ 'speech'
    (b) konuş$\langle$V$\rangle\langle$vn:inf$\rangle\langle$N$\rangle$ infinitive 'to speak', e.g., as in *konuşmamızı isemiyorlar* 'The do not want us *to speak*'
    (c) konuş$\langle$V$\rangle\langle$neg$\rangle\langle$imp$\rangle\langle$2s$\rangle$ 'speak-NEG-IMP = don't talk'

5. different affixes surfacing the same way, *evin* can be

    (a) ev-(n)In 'ev$\langle$N$\rangle\langle$gen$\rangle$ =of the house'
    (b) ev-(I)n 'ev$\langle$N$\rangle\langle$p2s$\rangle$ =your house'

6. The same surface suffix has multiple functions. For example, the word *doktorlar* can be,

    (a) doktor$\langle$N$\rangle\langle$pl$\rangle$ 'doctors'
    (b) doktor$\langle$N$\rangle\langle$0$\rangle\langle$V$\rangle\langle$cpl:pres$\rangle\langle$3p$\rangle$ 'they are doctors'

7. The suffix -(s)I that marks third person singular possessive and the null suffix that marks third person singular subject–predicate agreement may also have third person plural readings. For example,

    (a) The word *ev-i* can both mean 'his/her house' (ev$\langle$N$\rangle\langle$p3s$\rangle$) as well as 'their house' (ev$\langle$N$\rangle\langle$p3p$\rangle$).
    (b) A verb like *okudu* 'read-PAST' with no overt agreement marker may agree with a third person singular or plural subject. Hence, it is analyzed with both singular ('he/she read-PAST' oku$\langle$V$\rangle\langle$past$\rangle\langle$3s$\rangle$) and plural ('he/she read-PAST' oku$\langle$V$\rangle\langle$past$\rangle\langle$3p$\rangle$) third person agreement markers.

As a result, any predicate with a null agreement will have two analyses one with $\langle$3s$\rangle$ and the other with $\langle$3p$\rangle$ agreement tags. Similarly any noun with suffix -(s)I will have two analyses, one with $\langle$p3s$\rangle$ and the other with $\langle$p3p$\rangle$. These analyses will be multiplied with $\langle$ncomp$\rangle$ if the optional noun compound head marker is enabled during the compile time.

8. Some suffixes are not realized on the surface in the neighborhood of some other suffixes. These are generally, but not always, the suffixes having the same or similar surface forms. For example, *evleri* (the example in item 3) may be analyzed as

    (a) ev$\langle$N$\rangle\langle$p3p$\rangle$ as in *Annem ve babamın evleri İstanbul'da* 'My parents' house is in Istanbul'
    (b) ev$\langle$N$\rangle\langle$pl$\rangle\langle$p3p$\rangle$ as in *Annem ve babamın bütün evleri deniz manzaralı* 'All houses of my parents have a see view'.

since in case of $\langle$pl$\rangle$ (-lAr) and $\langle$p3p$\rangle$ (-lArI) are combined, the plural suffix -lAr does not realized on the surface.[14]

This particular source causes an extremely large number of ambiguous analyses because the multi functional suffix -(s)I is omitted in case it precedes (or follows) another -(s)I, but also a -lArI, -lI, -lIk, -slz, -CI or -CIk. Since some of these suffixes may

---

[14]One can also explain this as $\langle$p3p$\rangle$ being realized as -I in this particular context.

follow each other, and -(s)I itself has multiple functions, a word like *bağım-sız-lık-çı-lığ-ı-nı* causes a combinatorial expansion of ambiguous analyses because of the fact that at every suffix boundary marked with a dash in the example there may be a -(s)I suffix being deleted. This is further amplified by the fact that -(s)I may express ⟨ncomp⟩ or ⟨p3s⟩ and any of the resulting words may also have a null suffix expressing third person singular or plural agreement on a nominal predicate.[15] Most of these analyses will be semantically not plausible. However, there is no clear way of ruling them out at the analysis stage. The following illustrates the problem with a more tangible example, using the word *arabasız* which can be analyzed as one of the following (and more).

(c) araba⟨N⟩⟨siz⟩⟨Adj⟩ 'without a car'

(d) araba⟨N⟩⟨p3s⟩⟨siz⟩⟨Adv⟩ 'without a his/her car'

(e) araba⟨N⟩⟨ncomp⟩⟨siz⟩⟨Adv⟩, e.g., in *at arabasız* 'without a horse carriage'

(f) araba⟨N⟩⟨ncomp⟩⟨p3s⟩⟨siz⟩⟨Adv⟩, e.g., in *at arabasız* 'without his/her horse carriage'

Besides the ambiguity described above, overgeneration is another problem that one faces when the FST is used for generating surface forms. Unlike analysis, generation is almost always deterministic in Turkish. Nevertheless, there are a few cases where TRmorph produces multiple surface strings for a single analysis string. The following provides a (likely incomplete) list of cases where TRmorph is expected to overgenerate, i.e., either produce multiple (correct) surface strings for the same input, or produce incorrect surface strings in generation mode.

1. One of the clear cases where overgeneration occurs is the diminutive, ⟨dim⟩. The diminutive suffix in Turkish is one of -CIk, -cAk, -(I)cAk, -cAğız. TRmorph allows attaching any of these suffixes to any noun. This is unlikely to cause problems during the analysis. However, it will certainly produce incorrect surface forms.

---

[15]Most straightforward reading of the word is dative form of the noun phrase can roughly be translated as 'his/her state of being a supporter of independence'. With this root, The total number of analyses is 25560.

2. The ⟨p3s⟩ suffix -(s)I may also be used for marking third person plural possessive (⟨p3p⟩). For example *ev-i* in *Ali ve Ayşe'nin evi* 'The house of Ali and Ayşe' should be tagged as ⟨p3p⟩. On the other hand, the suffix -lArI is also used to express ⟨p3p⟩. As a result any analysis string with the symbol ⟨p3p⟩ will generate both surface options.

3. A similar case of overgeneration is with the null agreement suffix which should generally be tagged as ⟨3s⟩. However, such a predicate may also agree with a ⟨3p⟩ subject. Consequently, a null-agreement suffix on a predicate is tagged as both ⟨3s⟩ and ⟨3p⟩. Since ⟨3p⟩ can also be expressed with the suffix -lAr, a analysis string with ⟨3p⟩ also generates multiple surface forms.

4. Another known case of overgeneration is related to the relaxed analysis of alternative spellings or common misspellings. In the simplest case, every word will be generated once capitalized and once all lowercase. If 'all capitals' option is enabled, another surface form which is in all capital letters will be produced.

5. Similarly, if the analyzer is instructed to accept the proper noun suffixes without an apostrophe, in the generation mode the surface form with and without apostrophe will be included. As a result, some of the options may need to be tuned if the FST is to be used for generation.

6. Some symbols, like apostrophe have multiple representations in Unicode definition. As a result, any word that require an apostrophe will result in surface form for each alternative symbol.

7. After a small set of borrowings like *cami* 'mosque', the 's' in the suffix -(s)I is deleted according to official rules. However, this seems to be out of fashion in current use, and use of 's' (even in text) is more common that its deletion. Since TRmorph accepts both surface strings, this will cause generating multiple strings.

There are also a few other cases where some (sizable number of) speakers diverge from the canonical forms. An example is the redundant use of genitive suffix after a pronoun, before the suffix -(y)lA,

e.g., the surface form of '`sen`⟨`Prn:pers:1s`⟩⟨`ins`⟩' should be *sen-in-le* where the suffix -in is redundant. Some speakers tend not to use -in in such constructions. TRmorph accepts both use, hence the generation will be ambiguous.

8. Some borrowed words include a few vowels with circumflex, namely *â*, *û* and *î*. Except for a few words where use of circumflex helps disambiguation between different words, these vowels have been replaced by their non-circumflexed version in modern use. TRmorph allows this replacement even if the lexical form of the word should include a circumflex.[16] This also results in overgeneration, since any analysis string with a circumflexed vowel will have a surface form with and without circumflex.

6 Other tools

6.1 Stemming and lemmatization

In morphologically complex languages like Turkish, proper stemming requires analyzing the given word and stripping off the analysis symbols such that only the stem remains.

Although one can do this easily by filtering analyzer output, TRmorph includes a simple wrapper automaton for convenience. The automaton is defined in the file `stemmer.fst`. You need to type `make stemmer` to produce the binary `stem.fst`. This binary file can be used the way analyzer is used. Given a surface word, this automaton will produce the lexical form as the analysis string.

Optionally, one can keep the first tag, which is the syntactic category of the stem. Note that stemmer takes the lexical form as the 'stem', even if the lexical form has derivational suffixes immediately following the root form. Another compile time option related to stemmer causes the verbs to be suffixed with correct form of infinitive marker -mAk. This form of the verbs are what the dictionaries use as head words. Both options can be set in the file `options.h`.

---

[16]One can also allow circumflexed vowels to be used for their non-circumflexed counterparts in the lexicon. This is useful if one needs to analyze somewhat older text. Enabling this option will also cause overgeneration.

Note that ambiguity is less of a problem for the stemmer. However, in examples like *buna* discussed on page 13, there will be multiple stem forms produced (*bu*, *bun* and *buna* in this case).

6.2 Unknown word guesser

TRmorph includes a rudimentary guesser for guessing unknown words. To produce the automaton for this function, you should type `make guesser`, which would produce the file `guess.fst`. The usage of the automaton is again similar to the others. The surface strings of the FST is the (unknown) words, while analysis level is either the full analysis strings with possibly unknown root words that may lead to the surface form, or only the root word and its part of speech tag.

The guesser uses the same machinery as the analyzer, except the lexicon is replaced with a FSA that accepts a somewhat restricted set of strings as potential words. Since unknown words will likely include affixes, one may have a better chance of determining the root form of the word, and in most cases the class of the root word.

Depending on its application, the guesser be restricted further according to features of the words that can be coded into a finite state lexicon. For example, one may check whether the words fit into the syllable structure of the language, but this may miss the words of foreign origin, which are likely candidates for being unknown words. Currently only general restriction the guesser include the minimum and maximum root-word length that can be set in the file `options.h`.

The guesser may also be adjusted to return full analysis string(s) or only the root form followed by the POS tag. Again, these options can be set in `options.h`. Other customizations can be achieved by adjusting the file `guesser.lexc`.

The guesser is a standalone FST, to use it in combination with the analyzer, two automata can be combined with *priority union* such that guesser is only invoked if the analyzer fails. An example script that combines analyzer and guesser is provided with the name `anlyze_guess.xfst`. The resulting binary will be called `anlyze_guess.fst`. This binary is also built by default if the guesser is build by `make guesser`. By default the script puts two FST's without priority union to a single file. Hence, a command line like `flookup -a analyze_guess.fst` will simu-

late the priority union. Example command for building a single automaton combining the analyzer and the guesser with priority union is provided in `anlyze_guess.xfst`.

6.3 Morphological segmentation

Morphological segmentation is the task of finding morpheme boundaries on the surface strings. TRmorph distribution includes an automaton description for segmenting the words into their morphemes. To build the segmenter you need to type `make segmenter` and the resulting binary will be called `segment.fst`.

TRmorph marks the root and morpheme boundaries on the surface string to aid morpho-phonological rules. These boundaries are deleted from the surface string in the normal analyzer FST. The segmentation FST relies on this and the following trick for segmenting a given word to its surface morphemes: The given input string is first analyzed with the regular analyzer FST. Then the analysis strings are passed to a slightly modified FST in generation mode, which does not delete the boundary markers from the surface string.

It should be noted that the surface morpheme boundaries are not always determined uniquely. It is especially difficult to decide whether some buffer vowels or consonants belong to the morpheme preceding or following them. TRmorph consistently attaches these buffer letters to the morpheme that follow the boundary.

Because of the way it is implemented currently, the segmenter output needs to be post processed to obtain the desired result. The segmenter will produce multiple identical segmented strings, and there will also be some incorrect segmentations due to overgeneration discussed in Section 5. The output should be post-processed to remove multiple identical segmentations. The incorrect segmentations due to overgeneration can be eliminated by comparing the segmented string with the original one. An example post processing script is provided as `scripts/segment-filter.py`.

6.4 Hyphenation and syllabification

Hyphens in Turkish are inserted at the syllable boundaries. Because of the regular syllable structure and transparency of the orthography, this process does not require any dictionary lookup, or morphological analysis. Since the hy-phenation problem is easy to solve with a FST, a stand alone FST defined in xfst language included in the TRmorph distribution.

To build the hyphenation FST you need to type `make hyphenate` and the resulting binary will be called `hyphenate.fst`.

The surface string of the FST is Turkish words (or strings resembling words) and analysis string is the words where a hyphen '-' is inserted between the syllables, or at the points where one can insert a hyphen.

6.5 Analysis string converters

The distribution also includes a converter that does a best-effort conversion from the tagset generated by Oflazer (1994). It is a simple xfst script that takes Oflazer (1994) as surface, and generates TRmorph tags as the surface string. It was used/tested in in analysis mode (converting *to* TRmorph tags). However, it will produce multiple strings in the other direction, and the conversion *from* TRmorph is not tested, and may not work as expected. The conversion utility can be found as `converters/o2t.xfst`. The subdirectory contains a make file to build the binary file. The resulting binary will be called `o2t.fst`.

Another conversion utility that converts TRmorph output to a format similar to the output of MOR utility of CHILDES (MacWhinney and Snow 1985; MacWhinney and Snow 1990) is currently in development, will be made part of the standard TRmorph distribution soon.

6.6 Grapheme to phoneme conversion (g2p)

TRmorph includes an (experimental) grapheme to phoneme converter. The converter takes an analysis string and converts it to a phonemic representation in IPA. Turkish grapheme-to-phoneme conversion is (almost) trivial, the orthographic symbols map to phonetic symbols in a one-to-one fashion with simple context effects and a few exceptions. The conversion of orthographic symbols to the corresponding IPA symbols follows Göksel and Kerslake (2005). The tricky part is determining the location of stress.

In summary, the regular stress falls on the last syllable of a root form. With most suffixes, the stress moves

17

to the last syllable as new suffixes are added. For example the regular root *oku* 'read' followed by future TAM marker and first person singular agreement will result in IPA conversion okʰujɑdʒʰɑɣˈuɯm. Some suffixes contain non-stressable syllables. In those cases, stress moves to the syllable preceding the non-stressable syllable. For example, -yor in *okuyabiliyorum* will result in the IPA string okʰujɑbiɫˈijorum. Some root forms, notably palace names, compounds and some foreign borrowings show irregular stress. If a root is irregular, stress is on a non-final syllable, and stays at the same place with further suffixation with the exception of the negative suffix -mA. For example, for the irregular root *nitele* 'characterize/describe', the suffixation without a negative suffix will leave the stress on the first syllable (*niteliyorum* is pronounced nˈitʰeɫijorum), but negative marker move the stress to the syllable before itself (*nitelemiyourm* is pronounced nitʰeɫˈemijorum). Note that this will cause disambiguated surface forms to have ambiguous g2p conversions. For example the word *okuma* will be converted to both okʰumˈɑ 'oku⟨V⟩⟨vn:inf⟩⟨N⟩, meaning reading' and okʰˈumɑ 'oku⟨V⟩⟨neg⟩⟨imp⟩⟨2s⟩, don't read'.

TRmorph g2p component takes an analysis string and produces an IPA string in the generation mode. If the word does not have an exceptional stress marker in the lexicon, stress symbol is before the last vowel or the vowel that precedes a non-stressable syllable. The regular roots should not have any stress marking in the lexicon. If the exceptional stress is indicated in the lexicon, the affixes other than the negative marker do not change its location. Currently, the stress marker is placed before the vowel. The code for inserting syllable boundaries and the shifting the stress marker before the syllable is present in the file `g2p.xfst`, but it is not used by default.

As of this writing, g2p is is still experimental, and not tested throughly. Here are a list of things you should keep in mind while using this utility.

- The exceptional stress marking in the lexicon is far from complete. Especially, not all place names and compound forms are adequately marked.

- The foreign words, mainly proper names, that are written as in their origin are mispronounced (e.g., 'Coca Cola'). Currently there is no mechanism to mark these, or to specify an alternative pronunciation.

- Similar to the negative marker, Göksel and Kerslake (2005, p.30) list a number of other (polysyllabic) suffixes that cause some speakers to shift the stress to the preceding syllable regardless of the root form. TRmorph g2p does not produce multiple pronunciations for these suffixes.

- Since TRmorph only deals with the 'written words', the stress computation for compound words written separately, and clitics altering stress in the preceding word are not dealt with.

- Although TRmorph lexicon includes some exceptional stress marking, these are not checked well. The current markings are mainly based on Wiktionary definitions. Most of them are also checked automatically on online dictionary of Dil Derneği.

  It also seems some of these are difficult determine. For example the word *zatüre* 'pneumonia' is marked as stressed on the first syllable in the dictionary of Dil Derneği, but marked as finally stressed by Ergenç (2002). To add to the mix, my first intuition would be for medial stress.

7 Morphological disambiguator

TRmorph distribution currently includes a simple 'zero-context' disambiguator written in python. Work on disambiguation is still under way. The related files and this documentation will be updated with the developments on the disambiguator. The script for disambiguation is `scripts/disambiguate.py`. The script takes one token per line from its standard input, analyzes them using TRmorph, and prints a ranked list, by default with their scores (log probabilities). For details see the explanation on top of the script or the help text produced when run with the command line argument `--help`. The disambiguation script uses an external 'model file' An example model trained on the data set used by Yüret and Türe (2006) (originated from Hakkani-Tür et al. (2002)) can be found at `http://coltekin.net/cagri/trmorph/1M.m2`. With this setting, the expected accuracy of the disambiguator is about 85%. In other words, expect an error in every 7 tokens or so. The details on the method used for disambiguation and tools for training custom models will follow.

Acknowledgements

References

Akın, Ahmet Afşin and Mehmet Dündar Akın (2007). "Zemberek, an open source NLP framework for Turkic Languages". avialable from http://zemberek.googlecode.com/. URL: http://zemberek.googlecode.com/.

Beesley, Kenneth R. and Lauri Karttunen (2003). "Finite-state morphology: Xerox tools and techniques". In: *CSLI, Stanford*.

Çöltekin, Çağrı (2010). "A Freely Available Morphological Analyzer for Turkish". In: *Proceedings of the 7th International Conference on Language Resources and Evaluation (LREC 2010)*. Valetta, Malta, pp. 820–827.

Çöltekin, Çağrı (2014). "A set of open source tools for Turkish natural language processing". In: *Proceedings of the Ninth International Conference on Language Resources and Evaluation (LREC-2014)*. Reykjavik, Iceland: European Language Resources Association (ELRA).

Ergenç, İclal (2002). *Konuşma Dili ve Türkçenin Söyleyiş Sözlüğü*. Istanbul: Multilingual. ISBN: 9789756542064.

Göksel, Aslı and Celia Kerslake (2005). *Turkish: A Comprehensive Grammar*. London: Routledge.

Hakkani-Tür, Dilek Z., Kemal Oflazer, and Gökhan Tür (2002). "Statistical Morphological Disambiguation for Agglutinative Languages". In: *Computers and the Humanities* 36.4, pp. 381–410.

Hulden, Mans (2009). "Foma: a finite-state compiler and library". In: *Proceedings of the 12th Conference of the European Chapter of the Association for Computational Linguistics: Demonstrations Session*. Association for Computational Linguistics, pp. 29–32.

Lindén, Krister, Miikka Silfverberg, and Tommi Pirinen. "HFST Tools for Morphology–An Efficient Open-Source Package for Construction of Morphological Analyzers". In: *State of the Art in Computational Morphology*, pp. 28–47.

MacWhinney, Brian and Catherine Snow (1985). "The child language data exchange system". In: *Journal of Child Language* 12.2, pp. 271–269.

MacWhinney, Brian and Catherine Snow (1990). "The Child Language Data Exchange System: an update". In: *Journal of Child Language* 17.02, pp. 457–472.

Oflazer, Kemal (1994). "Two-level description of Turkish morphology". In: *Literary and Linguistic Computing* 9 (2).

Oflazer, Kemal and Gökhan Tür (1997). "Morphological Disambiguation by Voting Constraints". In: *Proceedings of the 35th Annual Meeting of the Association for Computational Linguistics*, pp. 222–229.

Schmid, Helmut (2005). "A programming language for finite state transducers". In: *Proceedings of the 5th International Workshop on Finite State Methods in Natural Language Processing (FSMNLP 2005)*. Helsinki, pp. 308–309.

Yüret, Deniz and Ferhan Türe (2006). "Learning morphological disambiguation rules for Turkish". In: *Proceedings of the Human Language Technology Conference of the North American Chapter of the Association of Computational Linguistics*. HLT-NAACL '06. New York, pp. 328–334. DOI: 10.3115/1220835.1220877.

---

[17]I started making this list 8 years after the first attempts at building TRmorph, so it is necessarily incomplete. I appreciate any reminders or corrections.

Table 7: Derivational morphemes analyzed by TRmorph. The column 'Derivation' lists the POS changes using a two letter symbols. The first letter is the original POS, and the second one is the POS after the suffixation. Here, N, J, A, M, V, I and O stand for noun, adjective, adverb, number, verb interjection and onomatopoeia, respectively.

Tag	Suffix	Derivation
⟨li⟩	-lI	NA NJ
⟨siz⟩	-sIz	NA NJ
⟨lik⟩	-lIk	NN JN AN
⟨dim⟩	-CIk	NN
	-cAk	
	-(I)cAk	
	-cAğIz	
⟨ci⟩	-CI	NN NJ
⟨arasi⟩	-arası	NJ
⟨imsi⟩	-(I)msI	NJ JJ
⟨imtrak⟩	-(I)mtrak	NJ JJ
⟨si⟩	-sI	NJ
⟨ca⟩	-CA	NA AA JJ MJ
⟨cana⟩	-CAnA	AA JA
⟨yici⟩	-(y)IcI	VJ
⟨cil⟩	-CIl	NJ
⟨gil⟩	-gil	NN
⟨lan⟩	-lAn	JV
⟨las⟩	-lAş	NV JV
⟨yis⟩	-yIş	VN
⟨esi⟩	-(y)AsI	VJ
⟨sal⟩	-sAl	NJ
⟨la⟩	-lA	NV JV OV IV
⟨dir⟩	-DIr	NA
⟨leri⟩	-lArI	AA

Index